

ADRIAN MATIAS

aadrian.mat@gmail.com | +34 654 435 598 | linkedin.com/in/adrianmatiasdiez | github.com/adrianmatias | x.com/adrian2mat |
adrianmatias.github.io Calle Miguel Cordero del Campillo 9, 24007 León – Spain

PROFILE

Data Scientist / ML Engineer with 12 years shipping end-to-end machine learning from problem framing to production. Deep roots in core DS — ad-tech ranking & user acquisition (Appodeal, Chartboost), fraud detection, NLP (RavenPack), recommender systems, Bayesian Bandits / RL, and large-scale experimentation on Spark / Databricks — with current work in GenAI: VET, PwC Advisory's LLM-driven vendor-evaluation platform (Azure OpenAI via Cavendish) — grounding raw vendor documentation into cited evidence, scoring engines, and contract-first pipelines. Held the correctness-gatekeeper seat: 15% of the codebase written, 52% of peer PRs reviewed. Approach driven by simplicity, first-principles thinking, and measurable business outcomes.

EXPERIENCE

Senior AI Engineer

Cavendish Professionals → PwC Advisory (long-term contract) | León, Spain – Remote | *May 2026 – Present*

Senior AI Engineer on the build of VET (Vendor Evaluation Tool), an LLM-driven platform PwC Advisory uses to evaluate vendors with contractually binding outcomes: ingest heterogeneous vendor submissions and raw documentation → ground extraction in cited evidence → score → compare. Ten-week engagement, four roles held in sequence — backend feature developer, contract & schema owner, Mode-1 workbook lane owner, and finally correctness gatekeeper & requirements arbiter. Wrote 15% of the codebase and reviewed 52% of every PR the team opened; the role drifted from writing features to deciding what counts as correct. Stack: Python, FastAPI, Pydantic, PostgreSQL, Azure OpenAI, Next.js, GitHub Actions; deployed on internal AKS.

Measured contribution — 10 weeks, 18 Jun – 28 Aug 2026

- **147 PRs authored, 116 merged** (79% of what opened; 139 in the main repo, 8 in CSX). Second-most of 6 by PR volume against a 366-PR staff counterpart — 139 vs 366 as a peer pair, not a hierarchy.
- **267 of 509 peer PRs reviewed** — 52% of everything the team opened; 486 review verdicts (39% of all review activity), 161 of them formal request-changes that blocked a merge until fixed. Review traffic with the staff counterpart is symmetric (212/204): a pair, not a seniority gradient.
- **150 of 994 commits on origin/main — 15% of the codebase**, plus 46 of 139 PRs taken to decision with no outside review: the backend carried correctness alone where the team had no second pair of eyes.
- **Median 3.0h open → merge, p90 3,775 lines, largest 20,834** — kept high-review-throughput and short cycle time at once (92% of merges under 24h) on a contract-and-plumbing-heavy codebase.
- **84 Jira items assigned / 74 raised** (61 completed, 41 owned by others after raising) across typed response envelopes, provenance gates, badge and rollback lane separation, the bus_reg contract tree, workbook parse / vendor match / averaging / rollup, and the CI defect-census and decision-ticket gates.
- **Enforcement built to outlive the engagement:** gate set, bus_reg contract tree, STE and thinness gates, and 161 request-change verdicts move correctness out of individual heads into gates that fail without the author — whoever inherits the repository inherits the enforcement.

What VET is

- AI-assisted vendor evaluation for contractual, binding decisions: raw submissions and vendor documentation grounded into cited, auditable evidence per requirement; heterogeneous data aggregated by domain experts; scoring paths validated for weighting defects (e.g. a headline vendor score that ignored framework dimension weights — diagnosed and fixed).
- Contract-first engineering: Pydantic as single source of truth, CI gates that fail on wire drift (camelCase/snake_case), schema-lane separation (badge / rollback), and a local end-to-end harness so assessment runs are reproducible outside the deployed stack.

Data Scientist → DS Technical Lead

Appodeal: appodeal.com | León, Spain – Remote | *Dec 2024 – Jan 2026*

Developed install prediction models for mobile user acquisition: cold-start, transformers, user embeddings, transfer learning, MLOps, production pipelines, parity resolution, experimentation, and monitoring. Progressed toward technical leadership based on ownership and strategic perspective.

Key Achievements

- Unified modeling across targets, advertisers, and platforms via attention/transformer architecture, increasing end KPI.
- Enabled seamless advertiser onboarding through novel cold-start techniques (custom PyTorch + transfer learning).
- Resolved parity mismatches and calibration gaps; productionized parity workflows with unit tests and decoupling in Databricks.
- Transfer-learning user embeddings achieving performance parity with orders-of-magnitude cost savings.

Senior Data Scientist

Quant AI Lab: quant.global | León, Spain – Remote | *Jul 2024 – Dec 2024*

Designed and delivered end-to-end ML solutions for international clients; reviewed technical scope, aligned with stakeholders, and supervised model productivization.

Key Contributions

- Led rapid turnaround of underperforming client ML system (utility fraud detection & ranking): diagnosed architectural issues, executed a surgical plan, restored accuracy/efficiency, and recovered client satisfaction.
- Discovered critical issues in data ingestion and feature-engineering pipelines.
- Led real-time credit card fraud classifier combining scalable ML with deep VAE. Open version: github.com/adrianmatias/ml-sandbox/tree/master/tx_class

Senior Data Scientist

Chartboost: chartboost.com | León, Spain – Remote | *Jun 2023 – Jun 2024*

Developed ad ranking and reinforcement learning systems for in-app game advertising; advanced DSP performance through large-scale experimentation and behavioral analysis.

Key Achievements

- Evolved and monitored Bayesian Bandit RL system with policy evaluation tools; designed tailored online/offline evaluation frameworks.
- Built Two-Towers neural ranking algorithm based on recommender systems.
- Post-mortem analysis of adversarial DSP dynamics and model feedback loops.

Machine Learning Engineer

RavenPack | León, Spain – Remote | *Mar 2022 – Jan 2023*

Researched, implemented, and maintained text analytics, NLU, and semantic search products on massive financial news data; productionized prototypes with emphasis on MLOps and reproducibility.

Key Achievements

- Few-shot search personalization: simplified architecture, reduced error 30%.
- Active learning PoC for learn-to-rank: projected 300% annotation efficiency gain.
- Novel long-text fingerprinting algorithm for duplicate content detection.

Earlier Roles

Numbrs *Aug 2020 – Oct 2021*

Extended Deep Learning / Active Learning for transaction processing; scaled classifier across countries with GDPR-compliant pipelines.

Stratio BD *Aug 2019 – Aug 2020*

Propensity/churn on Spark; BERT+graph email segmentation; multi-source weather DB (Elasticsearch/Flask).

Digilant *Jun 2015 – Apr 2017*

Redesigned Consumer Persona ML model; Game Theory attribution prototype; co-directed Master's thesis.

IEBS School *Oct 2019 – Oct 2021*

Webinars and masterclasses on real-time Big Data, AI, and ML applications.

DOCOMO DIGITAL *Apr 2017 – Aug 2019*

Led fraud detection (99.8% specificity, <50 ms) and default-payment filtering; mentored juniors.

2007–2015

Civil/tunnel engineering; university teaching: Math, Physics, Chemistry.

Freelance

Freelance Data Scientist / ML Engineer | *Mar 2022 – Jun 2025*

On-demand projects in semantic search, time series forecasting, and portfolio optimization for financial analysis.

SKILLS

Core Technologies

Python, FastAPI, Pydantic, PostgreSQL, PyTorch, Transformers, Azure OpenAI / Azure / AKS, AWS/S3, Databricks, PySpark, Pandas, SQL, Docker, MLflow, Git, GitHub Actions, AI-augmented tools (Ollama, llama.cpp, Hermes Agent)

Engineering Practices

Clean code, typed contracts as source of truth, CI drift gates, OOP & FP, TDD / local e2e harnesses, reproducible pipelines, Agile / XP, production ownership

Machine Learning & GenAI

LLMs, RAG, agentic systems, LangChain / LangGraph-style orchestration, document extraction & provenance, scoring engines, Deep Learning (attention, embeddings, transfer learning), Recommender Systems, NLP, RL / Bayesian Bandits, MLOps (parity, A/B, monitoring), Few-Shot / Active Learning, Ranking

Consulting & Soft Skills

Multi-client consulting delivery, stakeholder communication, correctness gatekeeping (486 review verdicts, 161 request-changes), mentoring, business impact focus, simplification / unification

SELECTED COURSES & CERTIFICATIONS

- Claude Code: A Highly Agentic Coding Assistant: Coursera (2026)
- ML Engineering for Production (MLOps) Specialization: Coursera (2021)
- Deep Learning with PyTorch: Siamese Networks (2022)
- Deep Learning Specialization: Coursera (2016)
- Structuring Machine Learning Projects: Coursera (2018)
- Apache Spark for Big Data & Scalable ML: Coursera (9.9/10, 2015)
- Functional Programming in Scala Specialization: Coursera (9.2/10, 2017)
- Bitcoin and Cryptocurrency Technologies: Coursera (9.5/10, 2017)
- Algorithms & Data Structures with Java (I, II): Coursera (8.5/10, 2014)
- Analytics Edge: Data Science & Predictive Models: MITx (9.3/10, 2014)
- Introduction to Programming & Data Science Using Python (I, II): MITx (8.7/10, 2014)

SELECTED PROJECTS

- **not the tao with adrian matias:** youtube.com/@notTheTaoWithAdrianMatias
Podcast exploring simplicity, consciousness, reality, machine learning, coding, sovereign AI.
- **RAG Blog System:** github.com/adrianmatias/ml-sandbox/tree/master/ragblog
Retrieval-augmented generation pipeline integrating LLMs, vector embeddings, and web crawling.
- **The Mirror:** <https://github.com/adrianmatias/the-mirror>
Knowledge repository based on llm-wiki for personalized assistant and insight for agentic AI.
- **Credit Card Transaction Fraud Classifier:** github.com/adrianmatias/ml-sandbox/tree/master/tx_class
Real-time fraud detection API (classical + unsupervised signals). Deploy-ready: FastAPI, Docker, RayServe.

EDUCATION

- **Master's in Geological Engineering:** University of Oviedo (2010–2012)
- **Degree in Mining & Energy Engineering:** University of León (2012)
- **Bachelor in Mining Engineering:** University of León (2003–2007)

References available on request